

# 1\_pipeline\_preprocessing

July 3, 2025

[ ]:

[1]:

```
"""  
Complete Scikit-Learn Preprocessing Pipeline for ISBSG Data  
=====
```

*This module provides a comprehensive preprocessing pipeline that handles:*

- 1. Data loading and initial cleaning*
- 2. Column name standardization*
- 3. Missing value handling*
- 4. Semicolon-separated value processing*
- 5. One-hot encoding for categorical variables*
- 6. Multi-label binarization for multi-value columns*
- 7. Feature selection and filtering*
- 8. Data validation and export*

*Based on the preprocessing steps from the provided notebooks.*

```
"""
```

```
[1]: '\n    Complete Scikit-Learn Preprocessing Pipeline for ISBSG Data\n    =====\n\n    This module provides a comprehensive preprocessing pipeline that handles:\n    1. Data loading and initial cleaning\n    2. Column name standardization\n    3. Missing value handling\n    4. Semicolon-separated value processing\n    5. One-hot encoding for categorical variables\n    6. Multi-label binarization for multi-value columns\n    7. Feature selection and filtering\n    8. Data validation and export\n\n    Based on the preprocessing steps from the provided notebooks.\n    '
```

[2]:

```
# === Imports ===  
  
import pandas as pd  
import numpy as np  
import re  
from pathlib import Path  
from sklearn.base import BaseEstimator, TransformerMixin  
from sklearn.pipeline import Pipeline  
from sklearn.preprocessing import MultiLabelBinarizer, StandardScaler
```

```

from sklearn.compose import ColumnTransformer
import joblib
import os
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
import warnings
from collections import Counter, defaultdict
from typing import List, Dict, Any, Tuple
warnings.filterwarnings('ignore')

```

```

[3]: # Configuration
DATA_FOLDER = "../data"
CONFIG_FOLDER = "../config"
SAMPLE_FILE = "ISBSG2016R1_1_financial_industry_seed.xlsx"
FULL_FILE = "ISBSG2016R1_1_full_dataset.xlsx"
TARGET_COL = "project_prf_normalised_work_effort" # be careful about case_
↳sensitive

```

```

[4]: from collections import Counter
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import numpy as np

def analyze_high_cardinality_multivalue(df, column, separator=';'):
    """
    Analyze high-cardinality multi-value columns to choose best strategy
    """
    print(f"=== ANALYSIS FOR HIGH-CARDINALITY COLUMN: '{column}' ===\n")

    # Basic statistics
    non_null_data = df[column].dropna().astype(str)
    split_values = non_null_data.apply(lambda x: [v.strip() for v in x.
↳split(separator) if v.strip()])

    # Get all unique values
    all_values = []
    for values_list in split_values:
        all_values.extend(values_list)

    value_counts = Counter(all_values)
    unique_values = list(value_counts.keys())

    print(f"Total unique values: {len(unique_values)}")
    print(f"Total value occurrences: {len(all_values)}")
    print(f"Average values per row: {len(all_values) / len(split_values):.2f}")

```

```

# Show most common values
print(f"\nTop 15 most common values:")
for value, count in value_counts.most_common(15):
    percentage = (count / len(non_null_data)) * 100
    print(f"  '{value}': {count} times ({percentage:.1f}% of rows)")

# Show distribution of value frequencies
frequency_dist = Counter(value_counts.values())
print(f"\nFrequency distribution:")
for freq, count in sorted(frequency_dist.items(), reverse=True)[:10]:
    print(f"  {count} values appear {freq} time(s)")

# Values per row distribution
values_per_row = split_values.apply(len)
print(f"\nValues per row:")
print(f"  Min: {values_per_row.min()}")
print(f"  Max: {values_per_row.max()}")
print(f"  Mean: {values_per_row.mean():.2f}")
print(f"  Median: {values_per_row.median():.2f}")

return value_counts, unique_values

def handle_high_cardinality_multivalue(df, multi_value_columns, separator=';',
    ↪strategy='top_k', preserve_original=None, **kwargs):
    """
    Handle high-cardinality multi-value columns with various strategies

    Parameters:
    -----
    strategy options:
    - 'top_k': Keep only top K most frequent values (k=kwargs['k'])
    - 'frequency_threshold': Keep values that appear in at least X% of rows
    ↪(threshold=kwargs['threshold'])
    - 'tfidf': Use TF-IDF vectorization with dimensionality reduction
    ↪(n_components=kwargs['n_components'])
    - 'count_features': Simple counting features (count, unique_count,
    ↪most_common)
    - 'embedding': Create category embeddings (requires pre-trained embeddings)
    """

    df_processed = df.copy()
    new_columns_mapping = {}
    preserve_original = preserve_original or []

    for col in multi_value_columns:

```

```

    if col not in df.columns:
        continue

    print(f"\nProcessing high-cardinality column '{col}' with strategy_
↳ '{strategy}'...")

    # Clean and split values
    split_values = df[col].fillna('').astype(str).apply(
        lambda x: [val.strip() for val in x.split(separator) if val.strip()]
    )

    # Get value counts
    all_values = []
    for values_list in split_values:
        all_values.extend(values_list)
    value_counts = Counter(all_values)

    if strategy == 'top_k':
        k = kwargs.get('k', 20) # Default to top 20
        top_values = [val for val, count in value_counts.most_common(k)]

        new_col_names = []
        for value in top_values:
            new_col_name = f"{col}_top_{value}".replace(' ', '_').
↳ replace('-', '_')
            df_processed[new_col_name] = split_values.apply(lambda x: 1 if_
↳ value in x else 0)
            new_col_names.append(new_col_name)

        # Add "other" category for remaining values
        other_col_name = f"{col}_other"
        df_processed[other_col_name] = split_values.apply(
            lambda x: 1 if any(val not in top_values for val in x) else 0
        )
        new_col_names.append(other_col_name)

        new_columns_mapping[col] = new_col_names
        print(f" Created {len(new_col_names)} columns (top {k} + other)")

    elif strategy == 'frequency_threshold':
        threshold = kwargs.get('threshold', 0.05) # Default 5%
        min_occurrences = int(len(df) * threshold)

        frequent_values = [val for val, count in value_counts.items() if_
↳ count >= min_occurrences]

        new_col_names = []

```

```

        for value in frequent_values:
            new_col_name = f"{col}_freq_{value}".replace(' ', '_').
↪replace('-', '_')
            df_processed[new_col_name] = split_values.apply(lambda x: 1 if
↪value in x else 0)
            new_col_names.append(new_col_name)

        # Add rare category
        rare_col_name = f"{col}_rare"
        df_processed[rare_col_name] = split_values.apply(
            lambda x: 1 if any(val not in frequent_values for val in x)
↪else 0
        )
        new_col_names.append(rare_col_name)

        new_columns_mapping[col] = new_col_names
        print(f"    Created {len(new_col_names)} columns
↪({len(frequent_values)} frequent + rare)")

    elif strategy == 'count_features':
        # Create aggregate features instead of individual columns
        new_col_names = []

        # Total count of values
        count_col = f"{col}_count"
        df_processed[count_col] = split_values.apply(len)
        new_col_names.append(count_col)

        # Unique count (in case of duplicates)
        unique_count_col = f"{col}_unique_count"
        df_processed[unique_count_col] = split_values.apply(lambda x:
↪len(set(x)))
        new_col_names.append(unique_count_col)

        # Most common value in the dataset appears in this row
        most_common_value = value_counts.most_common(1)[0][0] if
↪value_counts else None
        if most_common_value:
            most_common_col = f"{col}_has_most_common"
            df_processed[most_common_col] = split_values.apply(lambda x: 1
↪if most_common_value in x else 0)
            new_col_names.append(most_common_col)

        # Average frequency of values in this row
        avg_freq_col = f"{col}_avg_frequency"
        df_processed[avg_freq_col] = split_values.apply(

```

```

        lambda x: np.mean([value_counts[val] for val in x]) if x else 0
    )
    new_col_names.append(avg_freq_col)

    new_columns_mapping[col] = new_col_names
    print(f" Created {len(new_col_names)} aggregate feature columns")

elif strategy == 'tfidf':
    n_components = kwargs.get('n_components', 10) # Default to 10
    ↪components

    # Convert to text format for TF-IDF
    text_data = split_values.apply(lambda x: ' '.join(x))

    # Apply TF-IDF
    tfidf = TfidfVectorizer(max_features=100, stop_words=None)
    tfidf_matrix = tfidf.fit_transform(text_data)

    # Reduce dimensionality
    pca = PCA(n_components=n_components)
    tfidf_reduced = pca.fit_transform(tfidf_matrix.toarray())

    # Create new columns
    new_col_names = []
    for i in range(n_components):
        new_col_name = f"{col}_tfidf_comp_{i+1}"
        df_processed[new_col_name] = tfidf_reduced[:, i]
        new_col_names.append(new_col_name)

    new_columns_mapping[col] = new_col_names
    print(f" Created {len(new_col_names)} TF-IDF component columns")
    print(f" Explained variance ratio: {pca.
    ↪explained_variance_ratio_}")

elif strategy == 'hierarchical':
    # Group similar values into higher-level categories
    # This requires domain knowledge - example implementation
    hierarchy = kwargs.get('hierarchy', {}) # Dictionary mapping
    ↪values to categories

    if not hierarchy:
        print(" Warning: No hierarchy provided for hierarchical
    ↪strategy")
        continue

    # Create columns for each high-level category
    categories = set(hierarchy.values())

```

```

        new_col_names = []

        for category in categories:
            category_values = [val for val, cat in hierarchy.items() if cat
↪ == category]
            new_col_name = f"{col}_category_{category}".replace(' ', '_')
            df_processed[new_col_name] = split_values.apply(
                lambda x: 1 if any(val in category_values for val in x)
↪ else 0
            )
            new_col_names.append(new_col_name)

        new_columns_mapping[col] = new_col_names
        print(f"    Created {len(new_col_names)} hierarchical category
↪ columns")

        # FIXED: Only remove original column if NOT in preserve list
        if col not in preserve_original:
            df_processed = df_processed.drop(columns=[col])
            print(f"    Removed original column '{col}'")
        else:
            print(f"    Preserved original column '{col}' (in exclude list)")

    return df_processed, new_columns_mapping

def recommend_strategy(df, column, separator=';'):
    """
    Recommend the best strategy based on data characteristics
    """
    value_counts, unique_values = analyze_high_cardinality_multivaluedf,
↪ column, separator)

    total_unique = len(unique_values)
    total_rows = len(df[column].dropna())

    print(f"\n=== STRATEGY RECOMMENDATIONS FOR '{column}' ===")

    if total_unique > 100:
        print("    VERY HIGH CARDINALITY (100+ unique values)")
        print("Recommended strategies:")
        print("1. 'count_features' - Create aggregate features (safest)")
        print("2. 'top_k' with k=15-25 - Keep only most important values")
        print("3. 'tfidf' with n_components=5-10 - If values have semantic
↪ meaning")

    elif total_unique > 50:

```

```

print(" HIGH CARDINALITY (50+ unique values)")
print("Recommended strategies:")
print("1. 'top_k' with k=20-30 - Keep most frequent values")
print("2. 'frequency_threshold' with threshold=0.02-0.05")
print("3. 'count_features' - If you want aggregate information")

else:
    print(" MODERATE CARDINALITY (<50 unique values)")
    print("Recommended strategies:")
    print("1. 'frequency_threshold' with threshold=0.01")
    print("2. 'top_k' with k=30-40")
    print("3. Binary encoding might be acceptable")

# Check frequency distribution
freq_values = list(value_counts.values())
if max(freq_values) / min(freq_values) > 100:
    print("\n HIGHLY SKEWED DISTRIBUTION detected")
    print(" Consider 'frequency_threshold' or 'top_k' strategies")

```

```

[5]: def validate_multivalue_processing(df_original, df_processed, original_column,
    ↪ new_columns, separator=';', strategy='top_k'):
    """
    Comprehensive validation of multi-value categorical processing

    Parameters:
    -----
    df_original : pd.DataFrame
        Original dataset before processing
    df_processed : pd.DataFrame
        Processed dataset after handling multi-value columns
    original_column : str
        Name of original multi-value column
    new_columns : list
        List of new column names created from the original column
    separator : str
        Separator used in original data
    strategy : str
        Strategy used for processing
    """

    print(f"=== VALIDATION REPORT FOR COLUMN '{original_column}' ===\n")

    # 1. BASIC CHECKS
    print("1. BASIC INTEGRITY CHECKS")
    print("-" * 40)

    # Check row count consistency

```



```

original_rows = len(df_original)
processed_rows = len(df_processed)
print(f" Row count: {original_rows} → {processed_rows} {' SAME' if
↪original_rows == processed_rows else ' DIFFERENT'}")

# Check if original column was removed
original_removed = original_column not in df_processed.columns
print(f" Original column removed: {' YES' if original_removed else ' '
↪NO'}")

# Check if new columns exist
new_cols_exist = all(col in df_processed.columns for col in new_columns)
print(f" New columns created: {' YES' if new_cols_exist else ' NO'}
↪({len(new_columns)} columns)")

if not new_cols_exist:
    missing_cols = [col for col in new_columns if col not in df_processed.
↪columns]
    print(f" Missing columns: {missing_cols}")
    return False

# 2. DATA CONSISTENCY CHECKS
print(f"\n2. DATA CONSISTENCY CHECKS")
print("-" * 40)

# Parse original data
original_data = df_original[original_column].fillna('').astype(str)
split_original = original_data.apply(lambda x: [v.strip() for v in x.
↪split(separator) if v.strip()])

# Get all unique values from original
all_original_values = set()
for values_list in split_original:
    all_original_values.update(values_list)
all_original_values = sorted([v for v in all_original_values if v and v !=
↪'nan'])

print(f"Original unique values: {len(all_original_values)}")

if strategy == 'top_k':
    # Validate top-k strategy
    validate_top_k_strategy(df_original, df_processed, original_column,
↪new_columns, separator)
elif strategy == 'count_features':
    validate_count_features_strategy(df_original, df_processed,
↪original_column, new_columns, separator)

```

```

    elif strategy == 'frequency_threshold':
        validate_frequency_threshold_strategy(df_original, df_processed,
        ↪original_column, new_columns, separator)

    # 3. SAMPLE VALIDATION
    print(f"\n3. SAMPLE-BY-SAMPLE VALIDATION")
    print("-" * 40)
    validate_sample_rows(df_original, df_processed, original_column,
    ↪new_columns, separator, n_samples=5)

    # 4. STATISTICAL VALIDATION
    print(f"\n4. STATISTICAL VALIDATION")
    print("-" * 40)
    validate_statistics(df_original, df_processed, original_column,
    ↪new_columns, separator)

    # 5. INFORMATION LOSS ASSESSMENT
    print(f"\n5. INFORMATION LOSS ASSESSMENT")
    print("-" * 40)
    assess_information_loss(df_original, df_processed, original_column,
    ↪new_columns, separator)

    return True

def validate_top_k_strategy(df_original, df_processed, original_column,
    ↪new_columns, separator, k=None):
    """Validate top-k strategy specifically"""

    # Parse original data
    original_data = df_original[original_column].fillna('').astype(str)
    split_original = original_data.apply(lambda x: [v.strip() for v in x.
    ↪split(separator) if v.strip()])

    # Get value counts
    all_values = []
    for values_list in split_original:
        all_values.extend(values_list)
    value_counts = Counter(all_values)

    # Determine k if not provided
    if k is None:
        # Exclude "other" column to determine k
        non_other_cols = [col for col in new_columns if not col.
        ↪endswith('_other')]
        k = len(non_other_cols)

```

```

top_k_values = [val for val, count in value_counts.most_common(k)]
print(f"Top {k} values: {top_k_values[:5]}{'...' if len(top_k_values) > 5}
↳ else '']")

# Check each top-k column
for col in new_columns:
    if col.endswith('_other'):
        # Validate "other" column
        validate_other_column(df_original, df_processed, original_column,
↳ col, top_k_values, separator)
    else:
        # Extract the value name from column name
        value_name = col.replace(f"{original_column}_top_", "").
↳ replace(f"{original_column}_", "")
        validate_binary_column(df_original, df_processed, original_column,
↳ col, value_name, separator)

def validate_binary_column(df_original, df_processed, original_column,
↳ new_column, value_name, separator):
    """Validate a single binary column"""

    # Parse original data
    original_data = df_original[original_column].fillna('').astype(str)
    split_original = original_data.apply(lambda x: [v.strip() for v in x.
↳ split(separator) if v.strip()])

    # Expected values: 1 if value_name in the list, 0 otherwise
    expected = split_original.apply(lambda x: 1 if value_name in x else 0)
    actual = df_processed[new_column]

    # Compare
    matches = (expected == actual).sum()
    total = len(expected)
    match_rate = matches / total * 100

    print(f"  '{new_column}': {matches}/{total} matches ({match_rate:.1f}%")

    if match_rate < 100:
        mismatches = df_original.loc[expected != actual, original_column].
↳ head(3)
        print(f"    Sample mismatches: {list(mismatches)}")

```

```

def validate_other_column(df_original, df_processed, original_column,
    ↪other_column, top_values, separator):
    """Validate the 'other' category column"""

    # Parse original data
    original_data = df_original[original_column].fillna('').astype(str)
    split_original = original_data.apply(lambda x: [v.strip() for v in x.
    ↪split(separator) if v.strip()])

    # Expected: 1 if any value is NOT in top_values, 0 if all values are in
    ↪top_values
    expected = split_original.apply(lambda x: 1 if any(val not in top_values
    ↪for val in x) else 0)
    actual = df_processed[other_column]

    matches = (expected == actual).sum()
    total = len(expected)
    match_rate = matches / total * 100

    print(f" '{other_column}': {matches}/{total} matches ({match_rate:.1f}%)")

def validate_count_features_strategy(df_original, df_processed,
    ↪original_column, new_columns, separator):
    """Validate count features strategy"""

    # Parse original data
    original_data = df_original[original_column].fillna('').astype(str)
    split_original = original_data.apply(lambda x: [v.strip() for v in x.
    ↪split(separator) if v.strip()])

    for col in new_columns:
        if col.endswith('_count'):
            # Validate total count
            expected = split_original.apply(len)
            actual = df_processed[col]
            matches = (expected == actual).sum()
            print(f" '{col}': {matches}/{len(expected)} matches ({matches/
    ↪len(expected)*100:.1f}%)")

            elif col.endswith('_unique_count'):
                # Validate unique count
                expected = split_original.apply(lambda x: len(set(x)))
                actual = df_processed[col]
                matches = (expected == actual).sum()

```

```

        print(f" '{col}': {matches}/{len(expected)} matches ({matches/
↳len(expected)*100:.1f}%)")

def validate_frequency_threshold_strategy(df_original, df_processed,
↳original_column, new_columns, separator):
    """Validate frequency threshold strategy"""

    # Parse original data
    original_data = df_original[original_column].fillna('').astype(str)
    split_original = original_data.apply(lambda x: [v.strip() for v in x.
↳split(separator) if v.strip()])

    # Get value counts
    all_values = []
    for values_list in split_original:
        all_values.extend(values_list)
    value_counts = Counter(all_values)

    for col in new_columns:
        if col.endswith('_rare'):
            # Validate rare column - similar to other column validation
            continue
        else:
            # Extract the value name from column name
            value_name = col.replace(f"{original_column}_freq_", "").
↳replace(f"{original_column}_", "")
            validate_binary_column(df_original, df_processed, original_column,
↳col, value_name, separator)

def validate_sample_rows(df_original, df_processed, original_column,
↳new_columns, separator, n_samples=5):
    """Manually validate a few sample rows"""

    print(f"Validating {n_samples} random samples:")

    # Get random sample indices
    sample_indices = np.random.choice(len(df_original), min(n_samples,
↳len(df_original)), replace=False)

    for i, idx in enumerate(sample_indices, 1):
        original_value = df_original.iloc[idx][original_column]
        if pd.isna(original_value):
            original_values = []
        else:

```

```

        original_values = [v.strip() for v in str(original_value).
↪split(separator) if v.strip()]

    print(f"\n Sample {i} (row {idx}):")
    print(f"     Original: '{original_value}'")
    print(f"     Parsed: {original_values}")

    # Check new columns for this row
    for col in new_columns[:5]: # Show first 5 columns only
        processed_value = df_processed.iloc[idx][col]
        print(f"     {col}: {processed_value}")

def validate_statistics(df_original, df_processed, original_column,
↪new_columns, separator):
    """Validate statistical properties"""

    # Parse original data
    original_data = df_original[original_column].fillna('').astype(str)
    split_original = original_data.apply(lambda x: [v.strip() for v in x.
↪split(separator) if v.strip()])

    # Original statistics
    values_per_row = split_original.apply(len)
    print(f"Original values per row - Mean: {values_per_row.mean():.2f}, Std:
↪{values_per_row.std():.2f}")

    # New data statistics
    if any('_count' in col for col in new_columns):
        count_col = [col for col in new_columns if col.endswith('_count')][0]
        new_counts = df_processed[count_col]
        print(f"Processed counts - Mean: {new_counts.mean():.2f}, Std:
↪{new_counts.std():.2f}")

    # They should match!
    correlation = np.corrcoef(values_per_row, new_counts)[0, 1]
    print(f"Correlation between original and processed counts: {correlation:
↪.4f}")

    # Check for any impossible values
    binary_cols = [col for col in new_columns if not col.endswith(('_count',
↪'_frequency', '_avg_frequency'))]
    for col in binary_cols:
        unique_vals = df_processed[col].unique()
        if not set(unique_vals).issubset({0, 1, np.nan}):
            print(f" Warning: Non-binary values in '{col}': {unique_vals}")

```

```

def assess_information_loss(df_original, df_processed, original_column,
    ↪ new_columns, separator):
    """Assess how much information was lost in the transformation"""

    # Parse original data
    original_data = df_original[original_column].fillna('').astype(str)
    split_original = original_data.apply(lambda x: [v.strip() for v in x.
    ↪ split(separator) if v.strip()])

    # Get all unique values
    all_original_values = set()
    for values_list in split_original:
        all_original_values.update(values_list)
    all_original_values = sorted([v for v in all_original_values if v and v !=
    ↪ 'nan'])

    # Count how many unique values are captured by new columns
    captured_values = set()
    for col in new_columns:
        if not col.endswith(('_other', '_count', '_unique_count', '_frequency',
    ↪ '_avg_frequency', '_rare')):
            # Extract value name from column name
            value_parts = col.replace(f"{original_column}_", "").
    ↪ replace("top_", "").replace("freq_", "")
            captured_values.add(value_parts)

    capture_rate = len(captured_values) / len(all_original_values) * 100 if
    ↪ all_original_values else 0
    print(f"Value capture rate: {len(captured_values)}/
    ↪ {len(all_original_values)} ({capture_rate:.1f}%")

    if len(all_original_values) - len(captured_values) > 0:
        lost_values = set(all_original_values) - captured_values
        print(f"Lost values (first 10): {list(lost_values)[:10]}")

    # Estimate row-level information preservation
    if any('_other' in col for col in new_columns):
        other_col = [col for col in new_columns if col.endswith('_other')][0]
        rows_with_other = df_processed[other_col].sum()
        print(f"Rows with 'other' values: {rows_with_other}/{len(df_processed)}
    ↪ ({rows_with_other/len(df_processed)*100:.1f}%")

def quick_validation_summary(df_original, df_processed, column_mapping):

```

```

"""Quick validation summary for all processed columns"""

print("=== QUICK VALIDATION SUMMARY ===\n")

for original_col, new_cols in column_mapping.items():
    print(f" {original_col} → {len(new_cols)} new columns")

    # Check for obvious issues
    issues = []

    for col in new_cols:
        if col not in df_processed.columns:
            issues.append(f"Missing column: {col}")
        else:
            # Check for unexpected values in binary columns
            if not col.endswith(('_count', '_frequency', '_avg_frequency')):
                unique_vals = set(df_processed[col].dropna().unique())
                if not unique_vals.issubset({0, 1, 0.0, 1.0}):
                    issues.append(f"Non-binary values in {col}:␣
↪{unique_vals}")

    if issues:
        print(f"    Issues: {issues}")
    else:
        print(f"    Looks good")

print(f"\nDataset size: {len(df_original)} → {len(df_processed)} rows")
print(f"Column count: {len(df_original.columns)} → {len(df_processed.
↪columns)}")

```

```

[6]: def add_missing_categories_from_full_dataset(
    sample_df,
    full_df,
    categorical_columns,
    samples_per_category=2,
    exclude_columns=None
):
    """
    Add missing categorical values to sample dataset by sampling from full␣
    ↪dataset

    Parameters:
    -----
    sample_df : pd.DataFrame
        Your limited sample dataset
    full_df : pd.DataFrame
        Your complete dataset

```



```

categorical_columns : list
    List of categorical column names
samples_per_category : int
    Number of examples to add for each missing category
exclude_columns : list
    Columns that should not get ANY new categories (even indirectly)

Returns:
-----
pd.DataFrame : Enhanced dataset with missing categories included
"""

print("Analyzing missing categories...")

# Apply exclusions if provided at this level
if exclude_columns:
    categorical_columns = [col for col in categorical_columns
                           if col not in exclude_columns]
    print(f"Excluded columns: {exclude_columns}")

# Find missing categories in sample compared to full dataset
missing_categories = {}
category_stats = {}

for col in categorical_columns:
    if col not in sample_df.columns or col not in full_df.columns:
        print(f"Warning: Column '{col}' not found in one of the datasets")
        continue

    full_categories = set(full_df[col].dropna().unique())
    sample_categories = set(sample_df[col].dropna().unique())
    missing = full_categories - sample_categories

    if missing:
        missing_categories[col] = missing
        category_stats[col] = {
            'total_in_full': len(full_categories),
            'in_sample': len(sample_categories),
            'missing_count': len(missing)
        }
        print(f"Column '{col}': Missing {len(missing)} out of {len(full_categories)} categories")
        print(f"    Missing categories: {list(missing)[:5]}{'...' if len(missing) > 5 else ''}")
    else:
        print(f"Column '{col}': All categories present in sample")

```

```

if not missing_categories:
    print("No missing categories found! Your sample already contains all_
↪category values.")
    return sample_df.copy()

# Collect additional rows for missing categories
additional_rows = []
rows_added_by_category = defaultdict(int)

for col, missing_vals in missing_categories.items():
    print(f"\nSampling for column '{col}'...")

    for val in missing_vals:
        # Find all rows in full dataset with this category value
        matching_rows = full_df[full_df[col] == val]

        if len(matching_rows) == 0:
            print(f" Warning: No rows found for {col}='{val}' in full_
↪dataset")
            continue

        # Sample requested number of rows (or all available if fewer)
        n_samples = min(samples_per_category, len(matching_rows))
        sampled_rows = matching_rows.sample(n=n_samples, random_state=42)

        # Filter out rows that would introduce new categories in excluded_
↪columns
        if exclude_columns:
            original_sample_size = len(sampled_rows)

            for exclude_col in exclude_columns:
                if exclude_col in sampled_rows.columns and exclude_col in_
↪sample_df.columns:
                    # Get existing categories in sample
                    existing_categories = set(sample_df[exclude_col].
↪dropna().unique())

                    # Only keep rows where excluded column has existing_
↪values or is null
                    mask = (sampled_rows[exclude_col].
↪isin(existing_categories) |
                            sampled_rows[exclude_col].isna())
                    sampled_rows = sampled_rows[mask]

            if len(sampled_rows) == 0:

```

```

        print(f"  Skipped '{val}': Would violate exclusion_
↳constraints")

        continue
    elif len(sampled_rows) < original_sample_size:
        print(f"  Filtered {original_sample_size -
↳len(sampled_rows)} rows to respect exclusions")

        additional_rows.append(sampled_rows)
        rows_added_by_category[f"{col}='{val}'"] = len(sampled_rows)
        print(f"  Added {len(sampled_rows)} rows for '{val}' (out of_
↳{len(matching_rows)} available)")

    # Combine all additional rows
    if additional_rows:
        df_additional = pd.concat(additional_rows, ignore_index=True)

        # Remove potential duplicates (in case same row satisfies multiple_
↳missing categories)
        initial_additional_count = len(df_additional)
        df_additional = df_additional.drop_duplicates()
        final_additional_count = len(df_additional)

        if initial_additional_count != final_additional_count:
            print(f"\nRemoved {initial_additional_count -
↳final_additional_count} duplicate rows")

        # Combine with original sample
        df_enhanced = pd.concat([sample_df, df_additional], ignore_index=True)

        print(f"\n=== SUMMARY ===")
        print(f"Original sample size: {len(sample_df)}")
        print(f"Additional rows added: {len(df_additional)}")
        print(f"Final dataset size: {len(df_enhanced)}")
        print(f"Size increase: {len(df_additional)/len(sample_df)*100:.1f}%")

        return df_enhanced

    else:
        print("No additional rows could be sampled")
        return sample_df.copy()

def verify_categories_coverage(df_before, df_after, categorical_columns):
    """
    Verify that the enhanced dataset now covers all categories
    """
    print("\n=== CATEGORY COVERAGE VERIFICATION ===")

```

```

for col in categorical_columns:
    if col not in df_before.columns:
        continue

    before_cats = set(df_before[col].dropna().unique())
    after_cats = set(df_after[col].dropna().unique())
    new_cats = after_cats - before_cats

    print(f"\nColumn '{col}':")
    print(f"  Before: {len(before_cats)} categories")
    print(f"  After: {len(after_cats)} categories")
    if new_cats:
        print(f"  New categories added: {list(new_cats)}")

```

```

[7]: def create_preserving_pipeline(target_col, max_categorical_cardinality,
    ↪excluded_columns=None):
    """
    Create a pipeline that preserves excluded columns
    """
    excluded_columns = excluded_columns or []

    return Pipeline([
        ('multi_value_encoder', PreservingMultiValueEncoder(
            max_cardinality=max_categorical_cardinality,
            preserve_columns=excluded_columns
        )),
        ('categorical_encoder', PreservingCategoricalEncoder(
            max_cardinality=max_categorical_cardinality,
            preserve_columns=excluded_columns
        )),
        ('column_fixer', ColumnNameFixer()),
        ('validator', DataValidator(target_col))
    ])

```

```

[8]: import pandas as pd
import numpy as np
import re
import os
import joblib
from datetime import datetime
from pathlib import Path
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import MultiLabelBinarizer

# === 1. DataLoader: Load data and check target column ===

```

```

class DataLoader(BaseEstimator, TransformerMixin):
    """
        Load and perform initial data validation whether the target col exists:
        - Handles both .xlsx and .csv.
        - Stores the original shape of the data.
        - Raises an error if the target column is missing.

    """

    def __init__(self, file_path,
    ↪target_col='project_prf_normalised_work_effort'):
        self.file_path = file_path
        self.target_col = target_col # This should be the standardized form
        self.original_shape = None
        self.original_target_col = None # Store what we actually found

    def fit(self, X=None, y=None):
        return self

    def _standardize_column_name(self, col_name):
        """Convert column name to standardized format"""
        return col_name.strip().lower().replace(' ', '_')

    def _find_target_column(self, df_columns):
        """
        Smart target column finder - handles various formats
        Returns the actual column name from the dataframe
        """
        target_standardized = self.target_col.lower().replace(' ', '_')

        # Try exact match first
        if self.target_col in df_columns:
            return self.target_col

        # Try standardized versions of all columns
        for col in df_columns:
            col_standardized = self._standardize_column_name(col)
            if col_standardized == target_standardized:
                return col

        # If still not found, look for partial matches (for debugging)
        similar_cols = []
        target_words = set(target_standardized.split('_'))
        for col in df_columns:
            col_words = set(self._standardize_column_name(col).split('_'))
            if len(target_words.intersection(col_words)) >= 2: # At least 2
    ↪words match

```

```

        similar_cols.append(col)

    return None, similar_cols

def transform(self, X=None):
    """Load data from file with smart column handling"""

    print(f>Loading data from: {self.file_path}")

    # Determine file type and load accordingly; support for Excel or CSV
    if self.file_path.endswith('.xlsx'):
        df = pd.read_excel(self.file_path)
    elif self.file_path.endswith('.csv'):
        df = pd.read_csv(self.file_path)
    else:
        raise ValueError("Unsupported file format. Use .xlsx or .csv")

    self.original_shape = df.shape
    print(f>Loaded data with shape: {df.shape}")

    # Smart target column finding
    result = self._find_target_column(df.columns)

    if isinstance(result, tuple): # Not found, got similar columns
        actual_col, similar_cols = result
        error_msg = f"Target column '{self.target_col}' not found in data."
        if similar_cols:
            error_msg += f" Similar columns found: {similar_cols}"
        else:
            error_msg += f" Available columns: {list(df.columns)}"
        raise ValueError(error_msg)
    else:
        actual_col = result

    # Store the original column name we found
    self.original_target_col = actual_col

    if actual_col != self.target_col:
        print(f"Target column found: '{actual_col}' -> will be standardized_
↳to '{self.target_col}'")

    return df

# === 2. ColumnNameStandardizer: Clean and standardize column names ===
class ColumnNameStandardizer(BaseEstimator, TransformerMixin):
    """

```

*Standardize column names for consistency (lowercase, underscores, ↪ removes odd chars):*

- Strips spaces, lowercases, replaces & with \_&, removes special chars.*
- Useful for later steps and compatibility with modeling libraries.)*

```

"""

def __init__(self, target_col=None, original_target_col=None):
    self.column_mapping = {}
    self.target_col = target_col
    self.original_target_col = original_target_col

def fit(self, X, y=None):
    return self

def _standardize_columns(self, columns):
    """Standardize column names"""
    return [col.strip().lower().replace(' ', '_') for col in columns]

def _clean_column_names(self, columns):
    """Clean column names for compatibility"""
    cleaned_cols = []
    for col in columns:
        # Replace ampersands with _& to match expected transformations
        col_clean = col.replace(' & ', '_&_')
        # Remove special characters except underscores and ampersands
        col_clean = re.sub(r'[\w\s&]', '', col_clean)
        # Replace spaces with underscores
        col_clean = col_clean.replace(' ', '_')
        cleaned_cols.append(col_clean)
    return cleaned_cols

def transform(self, X):
    """Apply column name standardization"""
    df = X.copy()

    # Store original column names
    original_columns = df.columns.tolist()

    # Apply standardization
    standardized_cols = self._standardize_columns(original_columns)
    cleaned_cols = self._clean_column_names(standardized_cols)

    # Special handling for target column
    if self.original_target_col and self.target_col:
        target_index = None
        try:

```

```

        target_index = original_columns.index(self.original_target_col)
        cleaned_cols[target_index] = self.target_col
        print(f"Target column '{self.original_target_col}' -> '{self.
→target_col}'")
    except ValueError:
        pass # Original target col not found, proceed normally

    # Create mapping
    self.column_mapping = dict(zip(original_columns, cleaned_cols))

    # Apply new column names
    df.columns = cleaned_cols

    # Report changes
    changed_cols = sum(1 for orig, new in self.column_mapping.items() if
→orig != new)
    print(f"Standardized {changed_cols} column names")

    return df

# === 3. MissingValueAnalyzer: Analyze and handle missing values ===
class MissingValueAnalyzer(BaseEstimator, TransformerMixin):
    """
    Analyze and handle missing values
    - Reports number of columns with >50% and >70% missing.
    - Drops columns with a high proportion of missing data, except those
→you want to keep.
    - Fills remaining missing values:
        - Categorical: Fills with "Missing".
        - Numeric: Fills with column median.
    """

    def __init__(self, high_missing_threshold=0.7, cols_to_keep=None):
        self.high_missing_threshold = high_missing_threshold
        self.cols_to_keep = cols_to_keep or []
        self.high_missing_cols = []
        self.missing_stats = {}

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        """Analyze and handle missing values"""
        df = X.copy()

        # Calculate missing percentages

```



```

missing_pct = df.isnull().mean()
self.missing_stats = missing_pct.sort_values(ascending=False)

print(f"\nMissing value analysis:")
print(f"Columns with >50% missing: {sum(missing_pct > 0.5)}")
print(f"Columns with >70% missing: {sum(missing_pct > self.
↳high_missing_threshold)}")

    # Identify high missing columns
    self.high_missing_cols = missing_pct[missing_pct > self.
↳high_missing_threshold].index.tolist()

    # Filter out columns we want to keep
    final_high_missing_cols = [col for col in self.high_missing_cols if col
↳not in self.cols_to_keep]

    print(f"Dropping {len(final_high_missing_cols)} columns with >{self.
↳high_missing_threshold*100}% missing values")

    # Drop high missing columns
    df_clean = df.drop(columns=final_high_missing_cols)

    # Fill remaining missing values in categorical columns
    cat_cols = df_clean.select_dtypes(include=['object', 'category']).
↳columns
    for col in cat_cols:
        df_clean[col] = df_clean[col].fillna('Missing')

    # Fill remaining missing values in numerical columns with median
    num_cols = df_clean.select_dtypes(include=['number']).columns
    for col in num_cols:
        if df_clean[col].isnull().sum() > 0:
            median_val = df_clean[col].median()
            df_clean[col] = df_clean[col].fillna(median_val)
            print(f"Filled {col} missing values with median: {median_val}")

    print(f>Data shape after missing value handling: {df_clean.shape}")
    return df_clean

# === 4. SemicolonProcessor: Process multi-value columns (semicolon-separated)
↳===
class SemicolonProcessor(BaseEstimator, TransformerMixin):
    """
        Process semicolon-separated values in columns (e.g., "Python; Java;
↳SQL")
        - Identifies columns with semicolons.

```

```

        - Cleans: lowercases, strips, deduplicates, sorts, optionally
        ↪ standardizes values (e.g., "stand alone" → "stand-alone").
        - Useful for multi-value categorical features.

    """

    def __init__(self, standardization_mapping=None):
        self.semicolon_cols = []
        self.standardization_mapping = standardization_mapping or {
            "scrum": "agile development",
            "file &/or print server": "file/print server",
        }

    def fit(self, X, y=None):
        return self

    def _clean_and_sort_semicolon(self, val, apply_standardization=False,
        ↪ mapping=None):
        """Clean, deduplicate, sort, and standardize semicolon-separated
        ↪ values"""
        if pd.isnull(val) or val == '':
            return val

        parts = [x.strip().lower() for x in str(val).split(';') if x.strip()]

        if apply_standardization and mapping is not None:
            parts = [mapping.get(part, part) for part in parts]

        unique_cleaned = sorted(set(parts))
        return '; '.join(unique_cleaned)

    def transform(self, X):
        """Process semicolon-separated columns"""
        df = X.copy()

        # Identify columns with semicolons
        self.semicolon_cols = [
            col for col in df.columns
            if df[col].dropna().astype(str).str.contains(';').any()
        ]

        print(f"Found {len(self.semicolon_cols)} columns with semicolons: {self.
        ↪ semicolon_cols}")

        # Process each semicolon column
        for col in self.semicolon_cols:
            # Apply mapping for specific columns

```

```

        apply_mapping = col in ['process_pmf_development_methodologies',
        ↪ 'tech_tf_server_roles']
        mapping = self.standardization_mapping if apply_mapping else None

        # Clean the column
        df[col] = df[col].apply(
            lambda x: self._clean_and_sort_semicolon(x,
        ↪ apply_standardization=apply_mapping, mapping=mapping)
        )

    return df

# === 5. MultiValueEncoder: Encode semicolon columns using MultiLabelBinarizer
↪ ===
class MultiValueEncoder(BaseEstimator, TransformerMixin):
    """
        Handle multi-value columns using MultiLabelBinarizer
        - Only processes columns with a manageable number of unique values
        ↪ (max_cardinality).
        - Each semicolon column becomes several binary columns (e.g.,
        ↪ "lang__python", "lang__java", ...).
    """

    def __init__(self, max_cardinality=10):
        # Ensure max_cardinality is always an integer
        self.max_cardinality = int(max_cardinality) if max_cardinality is not
        ↪ None else 10
        self.multi_value_cols = []
        self.mlb_transformers = {}

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        """Encode multi-value columns"""
        df = X.copy()

        # Identify semicolon columns (multi-value)
        semicolon_cols = [
            col for col in df.columns
            if df[col].dropna().astype(str).str.contains(';').any()
        ]

        # Filter for low cardinality multi-value columns
        self.multi_value_cols = []
        for col in semicolon_cols:

```

```

        # Get unique values across all entries
        all_values = set()
        for val in df[col].dropna().astype(str):
            values = [v.strip() for v in val.split(';') if v.strip()]
            all_values.update(values)

        # Check cardinality (max_cardinality is already an integer from
↪__init__)
        if len(all_values) <= self.max_cardinality:
            self.multi_value_cols.append(col)

        print(f"Encoding {len(self.multi_value_cols)} multi-value columns:
↪{self.multi_value_cols}")

    # Process each multi-value column
    for col in self.multi_value_cols:
        # Prepare data for MultiLabelBinarizer
        values = df[col].dropna().astype(str).apply(
            lambda x: [item.strip() for item in x.split(';') if item.
↪strip()]
        )

        # Handle empty values - fill with empty list for MultiLabelBinarizer
        if len(values) == 0:
            continue

        # Fit and transform
        mlb = MultiLabelBinarizer()

        # Convert to list of lists, handling NaN/empty cases
        values_list = []
        for idx in df.index:
            if idx in values.index and values[idx]:
                values_list.append(values[idx])
            else:
                values_list.append([]) # Empty list for missing values

        onehot = pd.DataFrame(
            mlb.fit_transform(values_list),
            columns=[f"{col}__{cat}" for cat in mlb.classes_],
            index=df.index
        )

        # Store transformer for later use
        self.mlb_transformers[col] = mlb

    # Join with main dataframe

```

```

        df = df.join(onehot, how='left')

        print(f"Encoded {col} into {len(mlb.classes_)} binary columns")

        # Remove original multi-value columns
        df = df.drop(columns=self.multi_value_cols)

    return df

class PreservingMultiValueEncoder(MultiValueEncoder):
    """
    Modified MultiValueEncoder that preserves specific columns
    """
    def __init__(self, max_cardinality=10, preserve_columns=None):
        super().__init__(max_cardinality)
        self.preserve_columns = preserve_columns or []

    def transform(self, X):
        """Encode multi-value columns but preserve specified columns"""
        df = X.copy()

        # Store preserved columns before processing
        preserved_data = {}
        for col in self.preserve_columns:
            if col in df.columns:
                preserved_data[col] = df[col].copy()
                print(f"[PreservingMultiValueEncoder] Preserving column_{col}")

        # Apply normal multi-value encoding
        df_processed = super().transform(df)

        # Add back preserved columns (they might have been processed/removed)
        for col, data in preserved_data.items():
            if col not in df_processed.columns:
                df_processed[col] = data
                print(f"[PreservingMultiValueEncoder] Restored column '{col}'")

        return df_processed

# === 6. CategoricalEncoder: One-hot encode regular categorical columns ===
class CategoricalEncoder(BaseEstimator, TransformerMixin):
    """
    Handle single-value categorical columns
    - Ignores semicolon columns.

```

- Only encodes columns with a number of categories `max_cardinality`  
 ↳ (to avoid high-dimensional explosion).

- Can drop the first category for each variable to avoid  
 ↳ multicollinearity.

```

"""

def __init__(self, max_cardinality=10, drop_first=True):
    self.max_cardinality = max_cardinality
    self.drop_first = drop_first
    self.categorical_cols = []

def fit(self, X, y=None):
    return self

def transform(self, X):
    """Encode categorical columns"""
    df = X.copy()

    # Identify categorical columns
    cat_cols = df.select_dtypes(include=['object', 'category']).columns.
    ↳ tolist()

    # Identify semicolon columns to exclude
    semicolon_cols = [
        col for col in df.columns
        if df[col].dropna().astype(str).str.contains(';').any()
    ]

    # Filter for low cardinality single-value categorical columns
    excluded_columns = getattr(self, 'preserve_columns', []) # Add this
    ↳ attribute
    self.categorical_cols = [
        col for col in cat_cols
        if (col not in semicolon_cols and
            df[col].nunique() <= self.max_cardinality and
            col not in excluded_columns) # Skip excluded columns
    ]

    print(f"One-hot encoding {len(self.categorical_cols)} categorical
    ↳ columns: {self.categorical_cols}")

    # Apply one-hot encoding
    if self.categorical_cols:
        df = pd.get_dummies(df, columns=self.categorical_cols,
        ↳ drop_first=self.drop_first)

```

```

        return df

# Before creating the final pipeline, modify the encoder classes:
class PreservingCategoricalEncoder(CategoricalEncoder):
    """
    Modified CategoricalEncoder that preserves specific columns
    """
    def __init__(self, max_cardinality=10, drop_first=True,
        ↪preserve_columns=None):
        super().__init__(max_cardinality, drop_first)
        self.preserve_columns = preserve_columns or []

    def transform(self, X):
        """Encode categorical columns but preserve specified columns"""
        df = X.copy()

        # Store preserved columns before processing
        preserved_data = {}
        for col in self.preserve_columns:
            if col in df.columns:
                preserved_data[col] = df[col].copy()
                print(f"[PreservingCategoricalEncoder] Preserving column_
        ↪'{col}']")

        # Remove preserved columns from categorical processing
        original_cat_cols = df.select_dtypes(include=['object', 'category']).
        ↪columns.tolist()
        semicolon_cols = [
            col for col in df.columns
            if df[col].dropna().astype(str).str.contains(';').any()
        ]

        # Filter out preserved columns from processing
        self.categorical_cols = [
            col for col in original_cat_cols
            if (col not in semicolon_cols and
                df[col].nunique() <= self.max_cardinality and
                col not in self.preserve_columns) # Don't process preserved_
        ↪columns
        ]

        print(f"[PreservingCategoricalEncoder] Will encode {len(self.
        ↪categorical_cols)} columns (excluding {len(self.preserve_columns)}_
        ↪preserved)")

        # Apply one-hot encoding to non-preserved columns only

```

```

        if self.categorical_cols:
            df_processed = pd.get_dummies(df, columns=self.categorical_cols,
↳ drop_first=self.drop_first)
        else:
            df_processed = df.copy()

        # Ensure preserved columns are still there
        for col, data in preserved_data.items():
            if col not in df_processed.columns:
                df_processed[col] = data
                print(f"[PreservingCategoricalEncoder] Restored column '{col}'")

    return df_processed

# === 7. ColumnNameFixer: Final column name cleanup for PyCaret etc ===
class ColumnNameFixer(BaseEstimator, TransformerMixin):
    """
        Fix column names for PyCaret compatibility (removes illegal characters,
↳ replaces spaces/ampersands, handles duplicates):
        - No duplicate column names after encoding.
        - Only alphanumeric and underscores.

    """

    def __init__(self):
        self.column_transformations = {}

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        """Fix problematic column names"""
        df = X.copy()
        original_cols = df.columns.tolist()
        fixed_columns = []
        seen_columns = set()

        for col in original_cols:
            # Replace spaces with underscores
            fixed_col = col.replace(' ', '_')
            # Replace ampersands
            fixed_col = fixed_col.replace('&', 'and')
            # Remove other problematic characters
            fixed_col = ''.join(c if c.isalnum() or c == '_' else '_' for c in
↳ fixed_col)

```



```

        # Remove multiple consecutive underscores
        fixed_col = re.sub('_', '_', fixed_col)
        # Remove leading/trailing underscores
        fixed_col = fixed_col.strip('_')

        # Handle duplicates
        base_col = fixed_col
        suffix = 1
        while fixed_col in seen_columns:
            fixed_col = f"{base_col}_{suffix}"
            suffix += 1

        seen_columns.add(fixed_col)
        fixed_columns.append(fixed_col)

    # Store transformations
    self.column_transformations = dict(zip(original_cols, fixed_columns))

    # Apply new column names
    df.columns = fixed_columns

    # Check for duplicates
    dup_check = [item for item, count in pd.Series(fixed_columns).
        ↪value_counts().items() if count > 1]
    if dup_check:
        print(f"WARNING: Found {len(dup_check)} duplicate column names:␣
        ↪{dup_check}")
    else:
        print("No duplicate column names after fixing")

    n_changed = sum(1 for old, new in self.column_transformations.items()␣
    ↪if old != new)
    print(f"Fixed {n_changed} column names for PyCaret compatibility")

    return df

# === 8. DataValidator: Final summary and checks ===
class DataValidator(BaseEstimator, TransformerMixin):
    """
        Validate final dataset
        - Shape, missing values, infinities.
        - Data types (numeric, categorical).
        - Stats on the target column (mean, std, min, max, missing).
        - Report issues if any.

    """

```

```

def __init__(self, target_col):
    self.target_col = target_col

def fit(self, X, y=None):
    return self

def transform(self, X):
    """Validate the processed dataset"""
    df = X.copy()

    print(f"\n=== Final Data Validation ===")
    print(f"Final shape: {df.shape}")
    print(f"Target column: {self.target_col}")

    # Check for missing values
    missing_count = df.isnull().sum().sum()
    print(f"Total missing values: {missing_count}")

    # Check for infinite values
    numeric_cols = df.select_dtypes(include=[np.number]).columns
    inf_count = np.isinf(df[numeric_cols].values).sum()
    print(f"Total infinite values: {inf_count}")

    # Data types summary
    print(f"\nData types:")
    print(f"    Numeric columns: {len(df.select_dtypes(include=[np.number]).columns)}")
    print(f"    Categorical columns: {len(df.select_dtypes(include=['object', 'category']).columns)}")

    # Target variable summary
    if self.target_col in df.columns:
        target_stats = df[self.target_col].describe()
        print(f"\nTarget variable '{self.target_col}' statistics:")
        print(f"    Mean: {target_stats['mean']:.2f}")
        print(f"    Std: {target_stats['std']:.2f}")
        print(f"    Min: {target_stats['min']:.2f}")
        print(f"    Max: {target_stats['max']:.2f}")
        print(f"    Missing: {df[self.target_col].isnull().sum()}")
    else:
        print(f"WARNING: Target column '{self.target_col}' not found!")

    return df

# === Pipeline creation function: returns the Scikit-learn pipeline ===
def create_isbgs_preprocessing_pipeline(
    target_col='project_prf_normalised_work_effort',

```

```

original_target_col=None,
high_missing_threshold=0.7,
cols_to_keep=None,
max_categorical_cardinality=10,
standardization_mapping=None
):
    """
    Create complete preprocessing pipeline with smart target column handling

    Parameters:
    -----
    target_col : str
        Name of target column
    original_target_col : str
        Original target column name found in data
    high_missing_threshold : float
        Threshold for dropping columns with high missing values
    cols_to_keep : list
        Columns to keep even if they have high missing values
    max_categorical_cardinality : int
        Maximum number of unique values for categorical encoding
    standardization_mapping : dict
        Custom mapping for standardizing semicolon-separated values

    Returns:
    -----
    sklearn.pipeline.Pipeline
        Complete preprocessing pipeline
    """

    if cols_to_keep is None:
        cols_to_keep = [
            'project_prf_case_tool_used',
            'process_pmf_prototyping_used',
            'tech_tf_client_roles',
            'tech_tf_type_of_server',
            'tech_tf_clientserver_description'
        ]

    # Ensure max_categorical_cardinality is an integer
    if not isinstance(max_categorical_cardinality, int):
        max_categorical_cardinality = 10
        print(f"Warning: max_categorical_cardinality was not an integer,
↳ defaulting to {max_categorical_cardinality}")

    pipeline = Pipeline([

```

```

        ('column_standardizer', ColumnNameStandardizer(target_col,
↳original_target_col)),
        ('missing_handler', MissingValueAnalyzer(
            high_missing_threshold=high_missing_threshold,
            cols_to_keep=cols_to_keep
        )),
        ('semicolon_processor',
↳SemicolonProcessor(standardization_mapping=standardization_mapping)),
        ('multi_value_encoder',
↳MultiValueEncoder(max_cardinality=max_categorical_cardinality)),
        ('categorical_encoder',
↳CategoricalEncoder(max_cardinality=max_categorical_cardinality)),
        ('column_fixer', ColumnNameFixer()),
        ('validator', DataValidator(target_col))
    ])

    return pipeline

# === Full workflow function: orchestrates loading, pipeline, and saving ===
def preprocess_isbsg_data(
    file_path,
    target_col='project_prf_normalised_work_effort', # Always use standardized
↳form
    output_dir='../data',
    save_intermediate=True,
    **pipeline_kwargs
):
    """
    Complete preprocessing workflow for ISBSG data: loads the data, runs
    the full preprocessing pipeline, saves processed data, pipeline
    object, and a metadata report to disk, and returns the processed
    DataFrame and metadata

    Parameters:
    -----
    file_path : str
        Path to input data file
    target_col : str
        Name of target column
    output_dir : str
        Directory to save processed data
    save_intermediate : bool
        Whether to save intermediate processing steps
    **pipeline_kwargs : dict
        Additional arguments for pipeline creation

    Returns:

```

```

-----
pandas.DataFrame
    Processed dataframe ready for modeling
dict
    Processing metadata and statistics
"""

# print pipeline header
print("="*60)
print("ISBSG Data Preprocessing Pipeline")
print("="*60)
print(f"Processing file: {file_path}")
print(f"Target column (standardized): {target_col}")
print(f"Timestamp: {datetime.now()}")

# Create output directory
os.makedirs(output_dir, exist_ok=True)

# Load data with smart column detection
loader = DataLoader(file_path, target_col)
df_raw = loader.transform(X = None)

# Create and fit preprocessing pipeline
pipeline = create_isbsg_preprocessing_pipeline(
    target_col=target_col,
    original_target_col=loader.original_target_col, # Pass the found
column name
    **pipeline_kwargs
)

# Apply preprocessing in order of ColumnNameStandardizer=>
MissingValueAnalyzer =>
# SemicolonProcessor=> MultiValueEncoder=> CategoricalEncoder =>
ColumnNameFixer

# Apply preprocessing
df_processed = pipeline.fit_transform(df_raw)

# Prepare metadata
metadata = {
    'original_shape': loader.original_shape,
    'processed_shape': df_processed.shape,
    'processing_timestamp': datetime.now().isoformat(),
    'target_column_standardized': target_col,
    'target_column_original': loader.original_target_col,
    'pipeline_steps': [step[0] for step in pipeline.steps]
}

```

```

# Save processed data
file_stem = Path(file_path).stem
output_path = os.path.join(output_dir, f"{file_stem}_preprocessed.csv")
df_processed.to_csv(output_path, index=False)
print(f"\nProcessed data saved to: {output_path}")

# Save pipeline
pipeline_path = os.path.join(output_dir,
↪f"{file_stem}_preprocessing_pipeline.pkl")
joblib.dump(pipeline, pipeline_path)
print(f"Pipeline saved to: {pipeline_path}")

# Save metadata
metadata_path = os.path.join(output_dir,
↪f"{file_stem}_preprocessing_metadata.txt")
with open(metadata_path, 'w') as f:
    f.write("ISBSG Data Preprocessing Metadata\n")
    f.write("="*40 + "\n")
    for key, value in metadata.items():
        f.write(f"{key}: {value}\n")

print(f"Metadata saved to: {metadata_path}")

# Print completion & return results
print("\n" + "="*60)
print("Preprocessing completed successfully!")
print("="*60)

return df_processed, metadata

```

```

[9]: def integrated_categorical_preprocessing(
    sample_file_path: str,
    full_file_path: str,
    target_col: str,
    output_dir: str,
    cols_to_keep: List[str] = None,
    high_card_columns: List[str] = None,
    process_high_cardinality: bool = True,
    max_categorical_cardinality: int = 10,
    samples_per_category: int = 3,
    standardization_mapping: Dict[str, str] = None,
    high_missing_threshold: float = 0.7,
    separator: str = ';',
    strategy: str = 'top_k',
    k: int = 20,
    exclude_from_enhancement: List[str] = None

```

```

) -> Tuple[pd.DataFrame, Dict[str, Any]]:
    """
    Integrated pipeline to:
    1. Load sample and full datasets
    2. Apply consistent preprocessing to both datasets before comparison
    2. Auto-detect categorical columns
    3. Handle high-cardinality multi-value columns
    4. Enhance sample with missing categories from full dataset
    5. Apply standardization and final preprocessing

    Parameters:
    -----
    exclude_from_enhancement : List[str]
        List of column names to exclude from getting additional categories from
        ↪ full dataset

    Returns:
    - Enhanced and processed DataFrame
    - Metadata about the processing steps
    """

    print("="*60)
    print("INTEGRATED CATEGORICAL PREPROCESSING PIPELINE")
    print("="*60)

    # Initialize exclude list if not provided
    if exclude_from_enhancement is None:
        exclude_from_enhancement = []

    # Step 1: Load datasets
    print("\n1. Loading datasets...")
    sample_df = pd.read_excel(sample_file_path)
    full_df = pd.read_excel(full_file_path)

    # Lowercase all column names in both DataFrames independently
    sample_df.columns = [col1.lower() for col1 in sample_df.columns]
    full_df.columns = [col2.lower() for col2 in full_df.columns]

    print(f"Sample dataset shape: {sample_df.shape}")
    print(f"Full dataset shape: {full_df.shape}")

    # Step 2: Create preprocessing pipeline (WITHOUT final validation)
    print("\n2. Creating preprocessing pipeline...")

    # Create a pipeline that stops before final validation
    initial_pipeline = Pipeline([

```

```

        ('column_standardizer', ColumnNameStandardizer(target_col, target_col.
↳lower()))),
        ('missing_handler', MissingValueAnalyzer(
            high_missing_threshold=high_missing_threshold,
            cols_to_keep=cols_to_keep or []
        )),
        ('semicolon_processor',
↳SemicolonProcessor(standardization_mapping=standardization_mapping)),
    ])

    # Step 3: Apply initial preprocessing to BOTH datasets
    print("\n3. Applying initial preprocessing to both datasets...")

    # Process sample dataset
    sample_df_preprocessed = initial_pipeline.fit_transform(sample_df)
    print(f"Sample after initial preprocessing: {sample_df_preprocessed.shape}")

    # Process full dataset with same pipeline
    full_df_preprocessed = initial_pipeline.transform(full_df) # Use
↳transform, not fit_transform
    print(f"Full dataset after initial preprocessing: {full_df_preprocessed.
↳shape}")

    # Step 4: Handle high-cardinality columns on PREPROCESSED datasets
    # Initialize col_mapping regardless of whether we process high-cardinality
↳columns
    col_mapping = {}
    if process_high_cardinality and high_card_columns:
        print("\n4. Processing high-cardinality multi-value columns...")

        if high_card_columns is None:
            high_card_columns = ['external_eef_organisation_type',
↳'project_prf_application_type']

        # Filter out excluded columns from high-cardinality processing
        high_card_columns_to_process = [col for col in high_card_columns
            if col not in (exclude_from_enhancement or [])]

        if not high_card_columns_to_process:
            print("All high-cardinality columns are in exclusion list -
↳skipping processing")
        else:
            # Process high-cardinality columns in both datasets
            for col in high_card_columns_to_process:
                if col in full_df_preprocessed.columns:
                    print(f"\nProcessing high-cardinality column: {col}")

```



```

        # Process full dataset
        full_df_preprocessed, temp_mapping =
↪handle_high_cardinality_multivalue(
            full_df_preprocessed,
            multi_value_columns=[col],
            separator=separator,
            strategy=strategy,
            k=k,
            preserve_original=exclude_from_enhancement
        )

        # Process sample dataset with same strategy
        sample_df_preprocessed, _ =
↪handle_high_cardinality_multivalue(
            sample_df_preprocessed,
            multi_value_columns=[col],
            separator=separator,
            strategy=strategy,
            k=k,
            preserve_original=exclude_from_enhancement
        )

        col_mapping.update(temp_mapping)
    else:
        print("\n4. Skipping high-cardinality processing (disabled or no
↪columns specified)")

    # Step 4.5: Expand exclude list to include all derived binary columns
    print("\n4.5. Expanding exclude list to include derived columns...")
    expanded_exclude_list = exclude_from_enhancement.copy()

    for pattern in exclude_from_enhancement:
        # Find all columns that start with the excluded pattern
        matching_cols = [col for col in sample_df_preprocessed.columns
                        if col.startswith(pattern)]
        expanded_exclude_list.extend(matching_cols)
        if matching_cols:
            print(f" Added derived columns for '{pattern}': {matching_cols}")

    # Remove duplicates
    expanded_exclude_list = list(set(expanded_exclude_list))
    print(f" Final expanded exclude list: {len(expanded_exclude_list)}
↪columns")

    # Step 5: NOW identify categorical columns from preprocessed datasets
    print("\n5. Identifying categorical columns from preprocessed datasets...")

```

```

categorical_columns = []
for col in sample_df_preprocessed.columns:
    if (sample_df_preprocessed[col].dtype == 'object' or
        sample_df_preprocessed[col].nunique() <=
↪max_categorical_cardinality):
        categorical_columns.append(col)

categorical_columns = [col.lower() for col in categorical_columns]
print(f"Detected categorical columns: {len(categorical_columns)} columns")

# Step 6: Enhanced category sampling with PROPER exclusions
print("\n6. Enhancing sample with missing categories from preprocessed full_
↪dataset...")
print(f"Excluding columns from enhancement: {exclude_from_enhancement}")

# SEPARATE enhancement exclusion from processing exclusion
columns_to_enhance = [col for col in categorical_columns
                      if col not in (exclude_from_enhancement or [])]

print(f"Columns that will be enhanced: {len(columns_to_enhance)} out of_
↪{len(categorical_columns)}")

# Add missing categories ONLY for non-excluded columns
enhanced_df = add_missing_categories_from_full_dataset(
    sample_df=sample_df_preprocessed,
    full_df=full_df_preprocessed,
    categorical_columns=columns_to_enhance, # Only enhance these
    samples_per_category=samples_per_category,
    exclude_columns=exclude_from_enhancement # But respect exclusions for_
↪side effects
)

print(f"Enhanced dataset shape: {enhanced_df.shape}")

# CRITICAL FIX: Ensure excluded columns are still in enhanced_df
excluded_columns = exclude_from_enhancement or []
missing_excluded = []

for exclude_col in excluded_columns:
    if exclude_col not in enhanced_df.columns:
        missing_excluded.append(exclude_col)
        print(f"WARNING: Excluded column '{exclude_col}' missing from_
↪enhanced_df")

if missing_excluded:
    print(f"\nERROR: {len(missing_excluded)} excluded columns are missing!")
    print("This should not happen. Checking original preprocessed data...")

```

```

    for col in missing_excluded:
        if col in sample_df_preprocessed.columns:
            print(f" '{col}' exists in sample_df_preprocessed - copying_
↳over")
            enhanced_df[col] = sample_df_preprocessed[col]
        else:
            print(f" '{col}' not found anywhere!")

    # Step 6.5: Restore excluded columns if they got lost during_
↳high-cardinality processing
    print("\n6.5. Restoring excluded columns that were lost during_
↳high-cardinality processing...")

    excluded_columns = exclude_from_enhancement or []
    missing_excluded = []

    for exclude_col in excluded_columns:
        if exclude_col not in enhanced_df.columns:
            missing_excluded.append(exclude_col)
            print(f" WARNING: Excluded column '{exclude_col}' missing from_
↳enhanced_df")

    if missing_excluded:
        print(f"\nRestoring {len(missing_excluded)} excluded columns from_
↳original data...")

        for col in missing_excluded:
            # Try to find the column in the original sample data before_
↳high-cardinality processing
            if col in sample_df.columns: # Use the very original sample_df
                print(f" Restoring '{col}' from original sample_df")
                enhanced_df[col] = sample_df[col].iloc[:len(enhanced_df)] #_
↳Match the length
            elif col in sample_df_preprocessed.columns:
                print(f" Restoring '{col}' from sample_df_preprocessed")
                enhanced_df[col] = sample_df_preprocessed[col].iloc[:
↳len(enhanced_df)]
            else:
                print(f" ERROR: '{col}' not found in any source data!")

        print(f"Enhanced dataset shape after restoration: {enhanced_df.shape}")
    else:
        print(" All excluded columns are already present")

```

```

# Step 7: Verify categories coverage (ALL columns, but note which were
↪excluded)
print("\n7. Verifying categories coverage...")
print("\n=== CATEGORY COVERAGE VERIFICATION ===")
excluded_columns = exclude_from_enhancement or []

for col in categorical_columns:
    if col not in sample_df_preprocessed.columns:
        continue

    before_cats = set(sample_df_preprocessed[col].dropna().unique())
    after_cats = set(enhanced_df[col].dropna().unique())
    new_cats = after_cats - before_cats

    print(f"\nColumn '{col}':")
    print(f"  Before: {len(before_cats)} categories")
    print(f"  After:  {len(after_cats)} categories")

    if col in excluded_columns:
        if new_cats:
            print(f"    EXCLUDED column gained {len(new_cats)} categories_
↪(side effect): {list(new_cats)[:5]}")
        else:
            print(f"    EXCLUDED column preserved (no new categories)")
    else:
        if new_cats:
            print(f"    Enhanced with {len(new_cats)} new categories:
↪{list(new_cats)[:5]}")
        else:
            print(f"    No enhancement needed")

# Step 8: Apply final preprocessing stages to ALL columns
print("\n8. Applying final preprocessing stages...")

# Create modified pipeline that preserves excluded columns
final_pipeline = create_preserving_pipeline(
    target_col=target_col,
    max_categorical_cardinality=max_categorical_cardinality,
    excluded_columns=exclude_from_enhancement
)

final_df = final_pipeline.fit_transform(enhanced_df)

# Step 9: Final validation and duplicate check
print("\n9. Final validation and duplicate check...")

# Check for any remaining duplicates after all processing

```

```

    final_duplicate_cols = final_df.columns[final_df.columns.duplicated()].
    tolist()
    if final_duplicate_cols:
        print(f"Warning: Found duplicate columns in final dataset:␣
    {final_duplicate_cols}")
        final_df = final_df.loc[:, ~final_df.columns.duplicated()]
        print("Removed final duplicate columns")

    print(f"Original sample shape: {sample_df.shape}")
    print(f"Final processed shape: {final_df.shape}")
    print(f"Columns added: {final_df.shape[1] - sample_df.shape[1]}")
    print(f"Rows added: {final_df.shape[0] - sample_df.shape[0]}")

    # Compile metadata
    metadata = {
        'original_sample_shape': sample_df.shape,
        'original_full_shape': full_df.shape,
        'final_shape': final_df.shape,
        'categorical_columns_detected': categorical_columns,
        'high_cardinality_columns_processed': high_card_columns,
        'column_mapping': col_mapping,
        'rows_added_from_full_dataset': final_df.shape[0] - sample_df.shape[0]
    }

    return final_df, metadata

def safe_preprocess_with_fallback(
    enhanced_df: pd.DataFrame,
    target_col: str,
    output_dir: str,
    cols_to_keep: List[str] = None,
    max_categorical_cardinality: int = 10,
    standardization_mapping: Dict[str, str] = None,
    high_missing_threshold: float = 0.7
) -> Tuple[pd.DataFrame, Dict[str, Any]]:
    """
    Safe preprocessing function that handles the file_path requirement
    """

    # Save enhanced dataset to temporary file
    temp_enhanced_path = os.path.join(output_dir, 'temp_enhanced_sample.xlsx')
    enhanced_df.to_excel(temp_enhanced_path, index=False)

    try:
        # Apply preprocessing using existing function
        final_df, preprocessing_metadata = preprocess_isbsg_data(
            file_path=temp_enhanced_path,

```

```

        target_col=target_col,
        output_dir=output_dir,
        cols_to_keep=cols_to_keep,
        max_categorical_cardinality=max_categorical_cardinality,
        standardization_mapping=standardization_mapping,
        high_missing_threshold=high_missing_threshold
    )

    return final_df, preprocessing_metadata

finally:
    # Clean up temporary file
    try:
        os.remove(temp_enhanced_path)
    except:
        print(f"Warning: Could not remove temporary file_
↪{temp_enhanced_path}")

    return enhanced_df, {'error': 'Preprocessing failed'}

```

```

[10]: import os

# Configuration constants (define these at module level)
#DATA_FOLDER = "../data" # Update this path as needed
#SAMPLE_FILE = "sample_data.xlsx" # Update this filename as needed
#FULL_FILE = "full_data.xlsx" # Update this filename as needed
#TARGET_COL = "project_prf_normalised_work_effort"

print(f"\nDATA_FOLDER = {DATA_FOLDER}, SAMPLE_FILE = {SAMPLE_FILE}, FULL_FILE =_
↪{FULL_FILE}, TARGET_COL = {TARGET_COL}")

# Main execution function
def main():
    """
    Main function to run the integrated pipeline
    """

    # Configuration
    sample_file_path = os.path.join(CONFIG_FOLDER, SAMPLE_FILE)
    full_file_path = os.path.join(DATA_FOLDER, FULL_FILE)
    FINANCE = "finance"

    # Columns to exclude (customize as needed)
    cols_to_exclude_add_category = [
        'external_eef_industry_sector',
        'external_eef_organisation_type',

```

```

        'project_prf_application_type',
    ]

    # Columns to keep (customize as needed)
    cols_to_keep = [
        'Project_PRF_CASE_Tool_Used',
        'Process_PMF_Prototyping_Used',
        'Tech_TF_Client_Roles',
        'Tech_TF_Type_of_Server',
        'Tech_TF_ClientServer_Description'
    ]

    # High-cardinality multi-value columns: will top_k strategy - Keep only top_
    ↪K most frequent categorical values)
    high_card_columns = [
        'external_eef_organisation_type',
        'project_prf_application_type'
    ]

    # Standardization rules
    standardization_map = {
        'stand alone': 'stand-alone',
        'client server': 'client-server',
        'mathematically intensive': 'mathematically-intensive',
        ↪'#mathematically intensive application': 'mathematically-intensive_
    ↪application',
        "file &/or print server": "file/print server",
    }

    try:
        # Run integrated pipeline
        final_df, metadata = integrated_categorical_preprocessing(
            sample_file_path=sample_file_path,
            full_file_path=full_file_path,
            target_col=TARGET_COL,
            output_dir=DATA_FOLDER,
            cols_to_keep=cols_to_keep,
            high_card_columns=high_card_columns,
            max_categorical_cardinality=10,
            samples_per_category=3,
            standardization_mapping=standardization_map,
            high_missing_threshold=0.7,
            separator=';',
            strategy='top_k',
            k=20,
            process_high_cardinality=False,
            exclude_from_enhancement=cols_to_exclude_add_category

```

```

    )

    # Save results
    if FINANCE in sample_file_path:
        output_path = os.path.join(DATA_FOLDER,
        ↪f"{FINANCE}_enhanced_sample_final.csv")
    else:
        output_path = os.path.join(DATA_FOLDER, 'enhanced_sample_final.csv')

    final_df.to_csv(output_path, index=False)

    # Check what columns are actually in final_df
    print(f"\n=== FINAL COLUMN CHECK ===")
    print(f"Total columns in final CSV: {len(final_df.columns)}")
    print(f"All columns: {list(final_df.columns)}")

    print(f"\n" + "="*60)
    print("PIPELINE COMPLETED SUCCESSFULLY!")
    print("="*60)
    print(f"Final dataset saved to: {output_path}")
    print(f"Final shape: {final_df.shape}")
    print(f"Ready for PyCaret setup!")

    # Print summary of changes
    print(f"\nSUMMARY:")
    print(f"- Original sample rows: {metadata['original_sample_shape'][0]}")
    print(f"- Rows added from full dataset:
    ↪{metadata['rows_added_from_full_dataset']}")
    print(f"- Final rows: {metadata['final_shape'][0]}")
    print(f"- Original columns: {metadata['original_sample_shape'][1]}")
    print(f"- Final columns: {metadata['final_shape'][1]}")

    return final_df, metadata

except Exception as e:
    print(f"Error in integrated pipeline: {e}")
    raise

```

```

DATA_FOLDER = ../data, SAMPLE_FILE = ISBSG2016R1_1_financial_industry_seed.xlsx,
FULL_FILE = ISBSG2016R1_1_full_dataset.xlsx, TARGET_COL =
project_prf_normalised_work_effort

```

```
[ ]:
```



```
[11]: # Run the main function when script is executed directly
if __name__ == "__main__":
    final_df, metadata = main()
```

# ===== INTEGRATED CATEGORICAL PREPROCESSING PIPELINE =====

1. Loading datasets...

Sample dataset shape: (939, 51)

Full dataset shape: (7518, 52)

2. Creating preprocessing pipeline...

3. Applying initial preprocessing to both datasets...

Standardized 26 column names

Missing value analysis:

Columns with >50% missing: 29

Columns with >70% missing: 26

Dropping 26 columns with >70.0% missing values

Filled project\_prf\_functional\_size missing values with median: 154.5

Filled project\_prf\_normalised\_work\_effort\_level\_1 missing values with median: 1550.0

Filled project\_prf\_normalised\_work\_effort missing values with median: 1652.0

Filled project\_prf\_normalised\_level\_1\_pdr\_ufp missing values with median: 11.45

Filled project\_prf\_normalised\_pdr\_ufp missing values with median: 11.9

Filled project\_prf\_speed\_of\_delivery missing values with median: 27.05

Filled project\_prf\_project\_elapsed\_time missing values with median: 6.0

Data shape after missing value handling: (939, 25)

Found 2 columns with semicolons: ['external\_eef\_organisation\_type', 'project\_prf\_application\_type']

Sample after initial preprocessing: (939, 25)

Standardized 27 column names

Missing value analysis:

Columns with >50% missing: 30

Columns with >70% missing: 24

Dropping 24 columns with >70.0% missing values

Filled project\_prf\_functional\_size missing values with median: 139.0

Filled project\_prf\_normalised\_work\_effort\_level\_1 missing values with median: 1593.0

Filled project\_prf\_normalised\_work\_effort missing values with median: 1699.0

Filled project\_prf\_normalised\_level\_1\_pdr\_ufp missing values with median: 11.2

Filled project\_prf\_normalised\_pdr\_ufp missing values with median: 11.6

Filled project\_prf\_speed\_of\_delivery missing values with median: 26.8

Filled project\_prf\_project\_elapsed\_time missing values with median: 6.0

Filled project\_prf\_max\_team\_size missing values with median: 7.0

Data shape after missing value handling: (7518, 28)  
Found 4 columns with semicolons: ['external\_eef\_organisation\_type',  
'project\_prf\_application\_group', 'project\_prf\_application\_type',  
'process\_pmf\_development\_methodologies']  
Full dataset after initial preprocessing: (7518, 28)

4. Skipping high-cardinality processing (disabled or no columns specified)

4.5. Expanding exclude list to include derived columns...  
Added derived columns for 'external\_eef\_industry\_sector':  
['external\_eef\_industry\_sector']  
Added derived columns for 'external\_eef\_organisation\_type':  
['external\_eef\_organisation\_type']  
Added derived columns for 'project\_prf\_application\_type':  
['project\_prf\_application\_type']  
Final expanded exclude list: 3 columns

5. Identifying categorical columns from preprocessed datasets...  
Detected categorical columns: 14 columns

6. Enhancing sample with missing categories from preprocessed full dataset...  
Excluding columns from enhancement: ['external\_eef\_industry\_sector',  
'external\_eef\_organisation\_type', 'project\_prf\_application\_type']  
Columns that will be enhanced: 11 out of 14  
Analyzing missing categories...  
Excluded columns: ['external\_eef\_industry\_sector',  
'external\_eef\_organisation\_type', 'project\_prf\_application\_type']  
Column 'external\_eef\_data\_quality\_rating': All categories present in sample  
Column 'project\_prf\_application\_group': Missing 7 out of 7 categories  
Missing categories: ['real-time application', 'infrastructure software',  
'mathematically-intensive application', 'mathematically intensive application',  
'business application']...  
Column 'project\_prf\_development\_type': Missing 2 out of 7 categories  
Missing categories: ['Porting', 'Other']  
Column 'tech\_tf\_development\_platform': Missing 1 out of 7 categories  
Missing categories: ['Hand Held']  
Column 'tech\_tf\_language\_type': Missing 1 out of 7 categories  
Missing categories: ['APG']  
Column 'tech\_tf\_primary\_programming\_language': Missing 80 out of 129 categories  
Missing categories: ['ColdFusion', 'BPM', 'Azure', 'gcc', 'Adobe Flex']...  
Column 'project\_prf\_relative\_size': Missing 1 out of 10 categories  
Missing categories: ['XXXL']  
Column 'project\_prf\_case\_tool\_used': All categories present in sample  
Column 'tech\_tf\_architecture': Missing 2 out of 8 categories  
Missing categories: ['Multi-tier with web interface', 'Stand-alone']  
Column 'tech\_tf\_client\_server': All categories present in sample  
Column 'tech\_tf\_dbms\_used': All categories present in sample

Sampling for column 'project\_prf\_application\_group'...

- Skipped 'real-time application': Would violate exclusion constraints
- Skipped 'infrastructure software': Would violate exclusion constraints
- Skipped 'mathematically-intensive application': Would violate exclusion constraints
- Skipped 'mathematically intensive application': Would violate exclusion constraints
- Filtered 2 rows to respect exclusions
- Added 1 rows for 'business application' (out of 4655 available)
- Filtered 2 rows to respect exclusions
- Added 1 rows for 'missing' (out of 2311 available)
- Skipped 'business application; infrastructure software': Would violate exclusion constraints

Sampling for column 'project\_prf\_development\_type'...

- Skipped 'Porting': Would violate exclusion constraints
- Skipped 'Other': Would violate exclusion constraints

Sampling for column 'tech\_tf\_development\_platform'...

- Skipped 'Hand Held': Would violate exclusion constraints

Sampling for column 'tech\_tf\_language\_type'...

- Skipped 'APG': Would violate exclusion constraints

Sampling for column 'tech\_tf\_primary\_programming\_language'...

- Skipped 'ColdFusion': Would violate exclusion constraints
- Skipped 'BPM': Would violate exclusion constraints
- Skipped 'Azure': Would violate exclusion constraints
- Skipped 'gcc': Would violate exclusion constraints
- Skipped 'Adobe Flex': Would violate exclusion constraints
- Skipped 'Pega Workflows': Would violate exclusion constraints
- Skipped 'STAFFWARE': Would violate exclusion constraints
- Skipped 'IEF': Would violate exclusion constraints
- Skipped 'Huron/Object Star': Would violate exclusion constraints
- Skipped 'OutlookVBA': Would violate exclusion constraints
- Skipped 'VisualFoxPro': Would violate exclusion constraints
- Skipped 'PYTHON': Would violate exclusion constraints
- Skipped 'MS-Navision Properitory Language': Would violate exclusion constraints
- Skipped 'Siebel': Would violate exclusion constraints
- Skipped 'ADO.Net': Would violate exclusion constraints
- Skipped 'Delphi': Would violate exclusion constraints
- Skipped 'A:G': Would violate exclusion constraints
- Skipped 'RPL': Would violate exclusion constraints
- Skipped 'COGNOS': Would violate exclusion constraints
- Skipped 'BASIC': Would violate exclusion constraints
- Skipped 'APPS': Would violate exclusion constraints
- Skipped 'INGRES': Would violate exclusion constraints

Skipped 'BEA Weblogic': Would violate exclusion constraints  
 Skipped 'FORTRAN': Would violate exclusion constraints  
 Skipped 'ARBOR/BP': Would violate exclusion constraints  
 Skipped 'UNIFACE': Would violate exclusion constraints  
 Skipped 'Data base language': Would violate exclusion constraints  
 Skipped 'Periproducer': Would violate exclusion constraints  
 Skipped 'COOL:Gen': Would violate exclusion constraints  
 Skipped 'Centura': Would violate exclusion constraints  
 Skipped 'Datastage': Would violate exclusion constraints  
 Skipped 'Ada': Would violate exclusion constraints  
 Skipped 'Upfront': Would violate exclusion constraints  
 Skipped 'SAS': Would violate exclusion constraints  
 Skipped 'Doc1 Designer (Entorno visual)': Would violate exclusion constraints  
 Skipped 'Informatica PowerCenter': Would violate exclusion constraints  
 Skipped 'ACCEL': Would violate exclusion constraints  
 Skipped 'Visual Studio .Net': Would violate exclusion constraints  
 Skipped 'SLOGAN': Would violate exclusion constraints  
 Skipped 'Enablon': Would violate exclusion constraints  
 Skipped 'Caa': Would violate exclusion constraints  
 Skipped 'HPS': Would violate exclusion constraints  
 Skipped 'J2EE': Would violate exclusion constraints  
 Skipped 'MATLAB': Would violate exclusion constraints  
 Skipped 'LISP': Would violate exclusion constraints  
 Skipped 'BRE': Would violate exclusion constraints  
 Skipped 'XGML': Would violate exclusion constraints  
 Skipped 'BO': Would violate exclusion constraints  
 Skipped 'Express': Would violate exclusion constraints  
 Skipped 'AB INITIO': Would violate exclusion constraints  
 Skipped 'iPlanet Netscape Application Server': Would violate exclusion constraints  
 Skipped 'PERIPHONICS': Would violate exclusion constraints  
 Skipped 'EJB': Would violate exclusion constraints  
 Skipped 'PHP': Would violate exclusion constraints  
 Skipped 'TNSDL': Would violate exclusion constraints  
 Skipped 'IBM WTX': Would violate exclusion constraints  
 Skipped 'ADS/Online': Would violate exclusion constraints  
 Skipped 'Magic': Would violate exclusion constraints  
 Skipped 'Object oriented language': Would violate exclusion constraints  
 Skipped 'REXX': Would violate exclusion constraints  
 Skipped 'ABF': Would violate exclusion constraints  
 Skipped 'MANTIS': Would violate exclusion constraints  
 Skipped 'ASAP': Would violate exclusion constraints  
 Skipped 'NCR teradata scripting': Would violate exclusion constraints  
 Skipped 'Jdeveloper': Would violate exclusion constraints  
 Skipped 'C/AL': Would violate exclusion constraints  
 Skipped 'IIS': Would violate exclusion constraints  
 Skipped 'LEX': Would violate exclusion constraints  
 Skipped 'Perl': Would violate exclusion constraints

Skipped 'IDEAL': Would violate exclusion constraints  
Skipped 'Must Modeller': Would violate exclusion constraints  
Skipped 'Brightware proprietary': Would violate exclusion constraints  
Skipped 'SLEL': Would violate exclusion constraints  
Skipped 'PASCAL': Would violate exclusion constraints  
Skipped 'Spreadsheet': Would violate exclusion constraints  
Skipped 'Formspath': Would violate exclusion constraints  
Skipped 'DRIFT': Would violate exclusion constraints  
Skipped 'PowerPlay': Would violate exclusion constraints  
Skipped 'Mendix': Would violate exclusion constraints  
Skipped 'CICS': Would violate exclusion constraints

Sampling for column 'project\_prf\_relative\_size'...

Skipped 'XXXL': Would violate exclusion constraints

Sampling for column 'tech\_tf\_architecture'...

Skipped 'Multi-tier with web interface': Would violate exclusion constraints

Skipped 'Stand-alone': Would violate exclusion constraints

=== SUMMARY ===

Original sample size: 939

Additional rows added: 2

Final dataset size: 941

Size increase: 0.2%

Enhanced dataset shape: (941, 28)

6.5. Restoring excluded columns that were lost during high-cardinality processing...

All excluded columns are already present

7. Verifying categories coverage...

=== CATEGORY COVERAGE VERIFICATION ===

Column 'external\_eef\_data\_quality\_rating':

Before: 4 categories

After: 4 categories

No enhancement needed

Column 'external\_eef\_industry\_sector':

Before: 2 categories

After: 2 categories

EXCLUDED column preserved (no new categories)

Column 'external\_eef\_organisation\_type':

Before: 16 categories

After: 16 categories

EXCLUDED column preserved (no new categories)

Column 'project\_prf\_application\_group':  
Before: 5 categories  
After: 7 categories  
Enhanced with 2 new categories: ['missing', 'business application']

Column 'project\_prf\_application\_type':  
Before: 51 categories  
After: 51 categories  
EXCLUDED column preserved (no new categories)

Column 'project\_prf\_development\_type':  
Before: 5 categories  
After: 5 categories  
No enhancement needed

Column 'tech\_tf\_development\_platform':  
Before: 6 categories  
After: 6 categories  
No enhancement needed

Column 'tech\_tf\_language\_type':  
Before: 6 categories  
After: 6 categories  
No enhancement needed

Column 'tech\_tf\_primary\_programming\_language':  
Before: 49 categories  
After: 49 categories  
No enhancement needed

Column 'project\_prf\_relative\_size':  
Before: 9 categories  
After: 9 categories  
No enhancement needed

Column 'project\_prf\_case\_tool\_used':  
Before: 4 categories  
After: 4 categories  
No enhancement needed

Column 'tech\_tf\_architecture':  
Before: 6 categories  
After: 6 categories  
No enhancement needed

Column 'tech\_tf\_client\_server':  
Before: 5 categories

After: 5 categories  
No enhancement needed

Column 'tech\_tf\_dbms\_used':

Before: 3 categories  
After: 3 categories  
No enhancement needed

8. Applying final preprocessing stages...

[PreservingMultiValueEncoder] Preserving column 'external\_eef\_industry\_sector'  
[PreservingMultiValueEncoder] Preserving column 'external\_eef\_organisation\_type'  
[PreservingMultiValueEncoder] Preserving column 'project\_prf\_application\_type'  
Encoding 0 multi-value columns: []  
[PreservingCategoricalEncoder] Preserving column 'external\_eef\_industry\_sector'  
[PreservingCategoricalEncoder] Preserving column  
'external\_eef\_organisation\_type'  
[PreservingCategoricalEncoder] Preserving column 'project\_prf\_application\_type'  
[PreservingCategoricalEncoder] Will encode 12 columns (excluding 3 preserved)  
No duplicate column names after fixing  
Fixed 12 column names for PyCaret compatibility

=== Final Data Validation ===

Final shape: (941, 62)  
Target column: project\_prf\_normalised\_work\_effort  
Total missing values: 939  
Total infinite values: 0

Data types:

Numeric columns: 12  
Categorical columns: 4

Target variable 'project\_prf\_normalised\_work\_effort' statistics:

Mean: 4139.71  
Std: 9330.86  
Min: 6.00  
Max: 134211.00  
Missing: 0

9. Final validation and duplicate check...

Original sample shape: (939, 51)  
Final processed shape: (941, 62)  
Columns added: 11  
Rows added: 2

\n=== FINAL COLUMN CHECK ===

Total columns in final CSV: 62

All columns: ['isbsg\_project\_id', 'project\_prf\_year\_of\_project',  
'external\_eef\_industry\_sector', 'external\_eef\_organisation\_type',  
'project\_prf\_application\_type', 'tech\_tf\_primary\_programming\_language',

```

'project_prf_functional_size', 'project_prf_normalised_work_effort_level_1',
'project_prf_normalised_work_effort', 'project_prf_normalised_level_1_pdr_ufp',
'project_prf_normalised_pdr_ufp', 'project_prf_speed_of_delivery',
'project_prf_project_elapsed_time', 'process_pmf_docs', 'tech_tf_tools_used',
'project_prf_max_team_size', 'external_eef_data_quality_rating_B',
'external_eef_data_quality_rating_C', 'external_eef_data_quality_rating_D',
'project_prf_application_group_Infrastructure_Software',
'project_prf_application_group_Mathematically_Intensive_Application',
'project_prf_application_group_Missing',
'project_prf_application_group_Real_Time_Application',
'project_prf_application_group_business_application',
'project_prf_application_group_missing',
'project_prf_development_type_New_Development',
'project_prf_development_type_Not_Defined', 'project_prf_development_type_POC',
'project_prf_development_type_Re_development',
'tech_tf_development_platform_MR', 'tech_tf_development_platform_Missing',
'tech_tf_development_platform_Multi', 'tech_tf_development_platform_PC',
'tech_tf_development_platform_Proprietary', 'tech_tf_language_type_3GL',
'tech_tf_language_type_4GL', 'tech_tf_language_type_5GL',
'tech_tf_language_type_ApG', 'tech_tf_language_type_Missing',
'project_prf_relative_size_M1', 'project_prf_relative_size_M2',
'project_prf_relative_size_Missing', 'project_prf_relative_size_S',
'project_prf_relative_size_XL', 'project_prf_relative_size_XS',
'project_prf_relative_size_XXL', 'project_prf_relative_size_XXS',
'project_prf_case_tool_used_Missing', 'project_prf_case_tool_used_No',
'project_prf_case_tool_used_Yes', 'tech_tf_architecture_Missing',
'tech_tf_architecture_Multi_tier',
'tech_tf_architecture_Multi_tier_Client_server',
'tech_tf_architecture_Multi_tier_with_web_public_interface',
'tech_tf_architecture_Stand_alone', 'tech_tf_client_server_Missing',
'tech_tf_client_server_No', 'tech_tf_client_server_Not_Applicable',
'tech_tf_client_server_Yes', 'tech_tf_dbms_used_No', 'tech_tf_dbms_used_Yes',
'project_prf_team_size_group_Missing']

```

```

=====
PIPELINE COMPLETED SUCCESSFULLY!
=====

```

```

Final dataset saved to: ../data/enhanced_sample_final.csv
Final shape: (941, 62)
Ready for PyCaret setup!

```

#### SUMMARY:

- Original sample rows: 939
- Rows added from full dataset: 2
- Final rows: 941
- Original columns: 51
- Final columns: 62



[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	
[ ]:	